

CSC-462 ARTIFICIAL INTELLIGENCE

LAB PROJECT

Manaal Waseem
FA18-BCE-074
BCE-7B

Activities:

Task 1

Code:

```
In [164]: #import Libraries
import pandas as pd
import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline

#import 'r2_score'
from sklearn.metrics import r2_score

#import 'train_test_split'
from sklearn.model_selection import train_test_split

#import SimpleImputer
from sklearn.impute import SimpleImputer

#import MinMaxScaler
from sklearn.preprocessing import MinMaxScaler

#import pipeline
from sklearn.pipeline import Pipeline

from sklearn.model_selection import learning_curve
```

```
In [165]: #Load data
housing= pd.read_csv("housing.csv")
```

```
In [166]: #A Look at Housing data structure
housing.head()
```

```
Out[166]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_proximity
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0	NEAR BAY
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0	NEAR BAY
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0	NEAR BAY

In [167]: housing.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  -
0   longitude            20640 non-null  float64
1   latitude             20640 non-null  float64
2   housing_median_age   20640 non-null  float64
3   total_rooms          20640 non-null  float64
4   total_bedrooms       20433 non-null   float64
5   population           20640 non-null  float64
6   households           20640 non-null  float64
7   median_income        20640 non-null  float64
8   median_house_value   20640 non-null  float64
9   ocean_proximity      20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

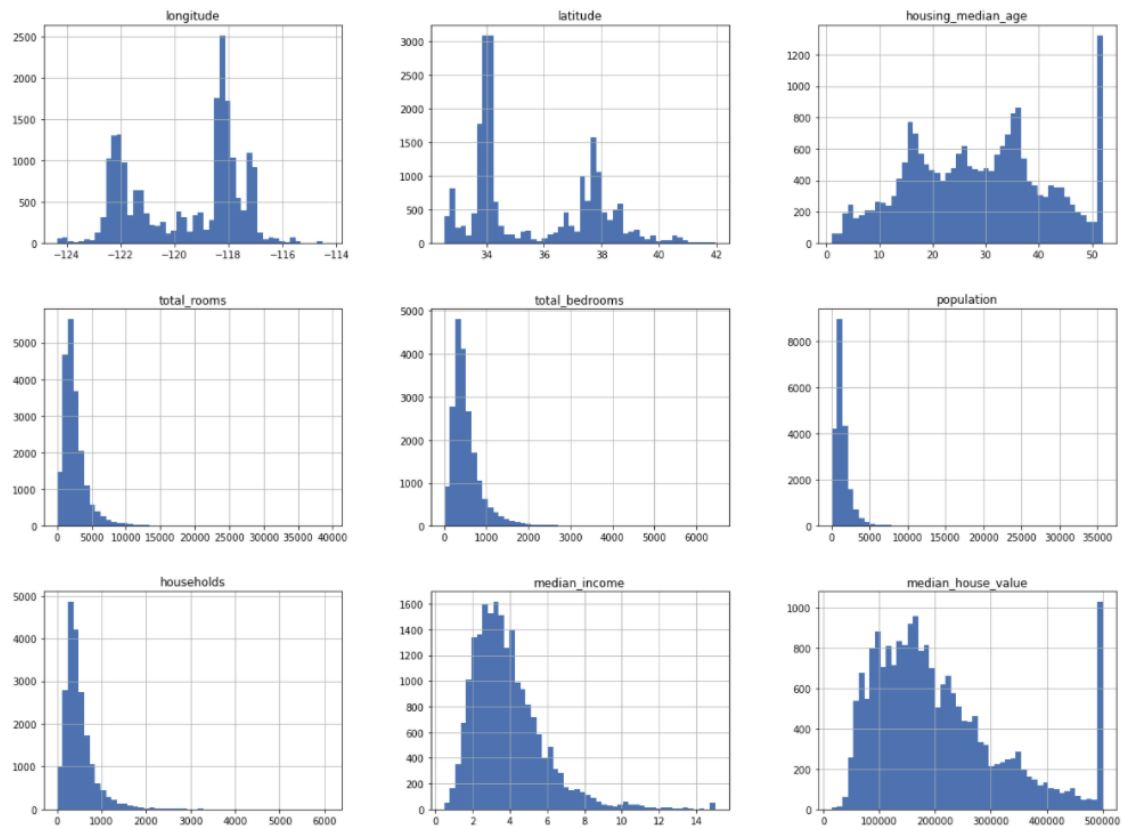
In [168]: housing.describe()

Out[168]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
count	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.000000	20640.000000	20640.000000	20640.000000
mean	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.476744	499.539680	3.870671	206855.816909
std	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.462122	382.329753	1.899822	115395.615874
min	-124.350000	32.540000	1.000000	2.000000	1.000000	3.000000	1.000000	0.499900	14999.000000
25%	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.000000	280.000000	2.563400	119600.000000
50%	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.000000	409.000000	3.534800	179700.000000
75%	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.000000	605.000000	4.743250	264725.000000
max	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.000000	6082.000000	15.000100	500001.000000


In [169]: *#drop ocean_proximity column as it is not required in our project*
housing = housing.drop("ocean_proximity", axis=1)

```
In [170]: #Data Exploration
housing.hist(bins=50, figsize=(20,15))
plt.show()
```



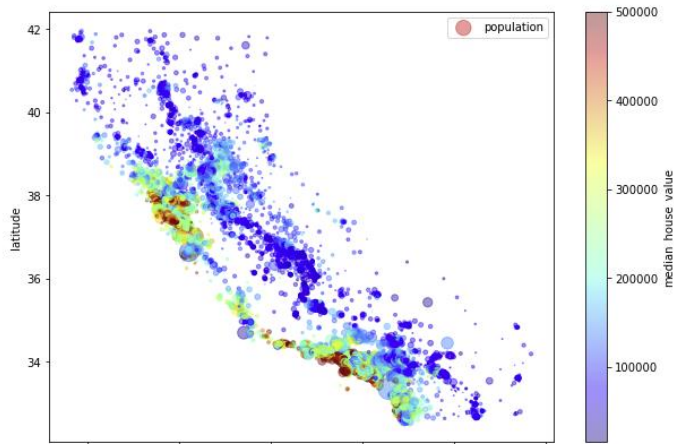
```
In [171]: #visualize the places where there is a high density of data points
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.1)
```

```
Out[171]: <AxesSubplot:xlabel='longitude', ylabel='latitude'>
```



```
In [172]: #a Look at the housing prices
housing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4,
s=housing["population"]/100, label="population", figsize=(10,7),
c="median_house_value", cmap=plt.get_cmap("jet"), colorbar=True,
)
plt.legend()
```

```
Out[172]: <matplotlib.legend.Legend at 0x22d08c35250>
```



```
In [173]: #Correlation matrix
corr_matrix = housing.corr()
corr_matrix["median_house_value"].sort_values(ascending=False)
```

```
Out[173]: median_house_value    1.000000
median_income      0.688075
total_rooms        0.134153
housing_median_age  0.105623
households         0.065843
total_bedrooms     0.049686
population         -0.024650
longitude          -0.045967
latitude           -0.144160
Name: median_house_value, dtype: float64
```

```
In [174]: housing = housing.dropna()
```

```
In [182]: housing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20433 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   longitude       20433 non-null  float64
1   latitude        20433 non-null  float64
2   housing_median_age  20433 non-null  float64
3   total_rooms     20433 non-null  float64
4   total_bedrooms  20433 non-null  float64
5   population      20433 non-null  float64
6   households      20433 non-null  float64
7   median_income   20433 non-null  float64
8   median_house_value 20433 non-null  float64
dtypes: float64(9)
memory usage: 1.6 MB
```

```
In [175]: #separate features X and target Y
Y = housing['median_house_value']
X = housing.drop('median_house_value', axis = 1)
```

```
In [176]: def performance_metric(y_true, y_predict):
score = r2_score(y_true, y_predict)
return score
```

```
In [177]: #split the data into training and testing subsets
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
print("Training and testing split was successful")
```

Training and testing split was successful

In [183]: X.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20433 entries, 0 to 20639
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   longitude              20433 non-null float64
1   latitude               20433 non-null float64
2   housing_median_age     20433 non-null float64
3   total_rooms            20433 non-null float64
4   total_bedrooms         20433 non-null float64
5   population             20433 non-null float64
6   households             20433 non-null float64
7   median_income          20433 non-null float64
dtypes: float64(8)
memory usage: 1.4 MB
```

In [207]: mean = X.mean(axis=0)
std = X.std(axis=0)

In [208]: X = X - mean
X = X / std

```
In [178]: from sklearn.base import BaseEstimator, TransformerMixin
rooms_ix, bedrooms_ix, population_ix, households_ix = 3, 4, 5, 6

class CombinedAttributesAdder(BaseEstimator, TransformerMixin):
    def __init__(self, add_bedrooms_per_room = True):
        self.add_bedrooms_per_room = add_bedrooms_per_room
    def fit(self, X, y=None):
        return self
    def transform(self, X, y=None):
        rooms_per_household = X[:, rooms_ix] / X[:, households_ix]
        population_per_household = X[:, population_ix] / X[:, households_ix]

        if self.add_bedrooms_per_room:
            bedrooms_per_room = X[:, bedrooms_ix] / X[:, rooms_ix]
            return np.c_[X, rooms_per_household, population_per_household, bedrooms_per_room]
        else:
            return np.c_[X, rooms_per_household, population_per_household]

attr_adder = CombinedAttributesAdder(add_bedrooms_per_room=False)
```

In []:

```
In [191]: #Pipeline Fill missing value, Combine attribute and Feature Normalization
# data_pipeline = Pipeline([
#     ('imputer', SimpleImputer(strategy="median")),
#     ('attrs_adder', CombinedAttributesAdder()),
#     ('min_max_scaler', MinMaxScaler()),
# ])

#X_pipeline = data_pipeline.fit_transform(X_train)
#scaler = MinMaxScaler()
#X_pipeline = scaler.fit_transform(X_train)
X_pipeline = X
```

```
In [215]: from sklearn.linear_model import LinearRegression
lin_reg = LinearRegression()
#lin_reg.fit(X_pipeline, Y_train)
lin_reg.fit(X, Y)
```

Out[215]: LinearRegression()

```
In [210]: test_df = pd.DataFrame({
    'longitude': [-119.85],
    'latitude': [37.48],
    'housing_median_age': [22],
    'total_rooms': [2850],
    'total_bedrooms': [500],
    'population': [1150],
    'households': [460],
    'median_income': [3.12]
})
```

```
In [211]: test_df
```

```
Out[211]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-119.85	37.48	22	2850	500	1150	460	3.12

```
In [212]: test_df = test_df - mean
test_df = test_df / std
```

```
In [213]: test_df
```

```
Out[213]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-0.139406	0.864456	-0.526779	0.097698	-0.089872	-0.242627	-0.103148	-0.395496

```
In [194]: # some_data = X_train.iloc[:5]
# some_labels = Y_train.iloc[:5]
# some_data_prepared = scalar.fit_transform(some_data)
# print("Predictions:", lin_reg.predict(some_data_prepared))
# print("Y:", list(some_labels))

some_data = X.iloc[:5]
some_labels = Y.iloc[:5]
some_data_prepared = scalar.fit_transform(some_data)
print("Predictions:", lin_reg.predict(some_data_prepared))
print("Y:", list(some_labels))

Predictions: [ -81100.51493603  165367.37007346  109819.06561593 -453833.36407158
  84713.18283089]
Y: [227600.0, 110400.0, 248100.0, 305600.0, 214600.0]
```

```
In [195]: from sklearn.metrics import mean_squared_error
housing_predictions = lin_reg.predict(X_pipeline)
lin_mse = mean_squared_error(Y_train, housing_predictions)
lin_rmse = np.sqrt(lin_mse)
lin_rmse
```

```
Out[195]: 69411.6554100667
```

```
In [196]: #R2 metric
r2=performance_metric(Y_train, housing_predictions)
r2
```

```
Out[196]: 0.6360185727313741
```

```
In [197]: train_sizes, train_scores, test_scores = learning_curve(estimator=LinearRegression(), X=X_train, y=Y_train,
    cv=10, train_sizes=np.linspace(0.1, 1.0, 10),
    n_jobs=1)
```


```
In [198]: # Learning Curve

size, train_scores, test_scores = learning_curve(lin, X_train, Y_train)
```

```
In [ ]: train_std = np.std(train_scores, axis=1)
test_std = np.std(test_scores, axis=1)
```

```
In [225]: plt.plot(size, train_std)
plt.plot(size, test_std)
plt.legend(['Train', 'Test'])
plt.title('Learning Curve')
plt.xlabel('Training Size')
plt.ylabel('Loss Score')
```

```
Out[225]: Text(0, 0.5, 'Loss Score')
```



```
In [200]: # values of the trained weights
print(lin_reg.coef_)
```

```
[ -420361.45533591  -398382.16450844   58909.95253161  -331860.3736203
  746410.9383576   -1311516.057383   253215.48155619   584260.59557271]
```

```
In [201]: #Test cases
X_test_pipeline = scalar.fit_transform(X_test)
final_predictions = lin_reg.predict(X_test_pipeline)
final_mse = mean_squared_error(Y_test, final_predictions)
final_rmse = np.sqrt(final_mse)
```

```
In [202]: # final_rmse
final_rmse
```

```
Out[202]: 92901.10019388044
```

```
In [203]: #R2 metric
r2=performance_metric(Y_test, final_predictions)
r2
```

```
Out[203]: 0.368884802895016
```


```
In [204]: #final training loss value
print('Your final training loss value rmse = {} R2 = {}'.format(final_rmse, r2))

Your final training loss value rmse = {} R2 = {} 92901.10019388044 0.368884802895016
```

```
In [224]: print(f'Predicted value for median house value: {lr.predict([[-119.85,37.48,22,2850,500,1150,400,3.12]])}')
print(f'Value of training weights: {lr.coef_}')
actual = np.array(list(Y_train))
predicted = np.array(list(pri))
print(mean_squared_error(actual,predicted,squared=False))

Predicted value for median house value: [102411.04824779]
Value of training weights: [-4.30776320e+04 -4.28059533e+04  1.13640741e+03 -8.26537227e+00
  1.19764264e+02 -3.81207981e+01  4.00312819e+01  4.01947682e+04]
69633.72002531342
```

Learning Curve



Weights

```
In [154]: # values of the trained weights
print(lin_reg.coef_)

[ -418268.88736635  -387070.03453979   60557.81471414   77206.23571651
   91860.60238544  -1469411.46051484   663328.64835277   618217.58800007
   495057.21138978    71127.52066782   292806.6403382 ]
```

Training Loss Value:

```
In [161]: #final training loss value
print('Your final training loss value rmse = {} R2 = {}'.format(final_rmse, r2))

Your final training loss value rmse = {} R2 = {} 74292.19350137576 0.5788085928215299
```

Predicted Value for Test Case:

```
Predicted value for median house value: [102411.04824779]
Value of training weights: [-4.30776320e+04 -4.28059533e+04  1.13640741e+03 -8.26537227e+00
  1.19764264e+02 -3.81207981e+01  4.00312819e+01  4.01947682e+04]
69633.72002531342
```


Task 2

Code:

```
In [3]: import tensorflow.keras as keras
import pandas as pd
import numpy as numpy
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [28]: df = pd.read_csv('housing.csv')
```

```
In [3]: df.shape
```

```
Out[3]: (20640, 10)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   longitude            20640 non-null  float64
1   latitude             20640 non-null  float64
2   housing_median_age   20640 non-null  float64
3   total_rooms          20640 non-null  float64
4   total_bedrooms       20433 non-null  float64
5   population           20640 non-null  float64
6   households           20640 non-null  float64
7   median_income        20640 non-null  float64
8   median_house_value   20640 non-null  float64
9   ocean_proximity      20640 non-null  object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
```

```
In [5]: df = df.dropna()
```

first 8 columns are features, 9th col is label and 10th is not used

```
In [6]: labels = df.median_house_value
```

```
In [7]: features = df.copy()
```

```
In [8]: features.drop(columns=['median_house_value', 'ocean_proximity'], inplace=True)
```

```
In [9]: features.head()
```

```
Out[9]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462

```
In [10]: features.shape
```

```
Out[10]: (20433, 8)
```

```
In [11]: labels.shape
```

```
Out[11]: (20433,)
```

Normalization

```
In [12]: mean = features.mean(axis=0)
std = features.std(axis=0)
```

```
In [13]: features = features - mean
features = features / std
```

```
In [14]: mean
```

```
Out[14]: longitude      -119.570689
latitude         35.633221
housing_median_age  28.633094
total_rooms      2636.504233
total_bedrooms   537.870553
population       1424.946949
households       499.433465
median_income     3.871162
dtype: float64
```

```
In [15]: features.head()
```

```
Out[15]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-1.327281	1.051692	0.982139	-0.803793	-0.970301	-0.973296	-0.976809	2.345106
1	-1.322290	1.042330	-0.606195	2.042080	1.348243	0.861318	1.670332	2.332575
2	-1.332272	1.037649	1.855723	-0.535176	-0.825541	-0.819749	-0.843406	1.782896
3	-1.337263	1.037649	1.855723	-0.623495	-0.718750	-0.765037	-0.733544	0.932947
4	-1.337263	1.037649	1.855723	-0.461959	-0.611959	-0.758860	-0.628914	-0.013143

```
In [16]: labels.head()
```

```
Out[16]: 0    452600.0
1    358500.0
2    352100.0
3    341300.0
4    342200.0
Name: median_house_value, dtype: float64
```

Neural Network Model

```
In [17]: model = keras.models.Sequential()
model.add(keras.layers.Dense(8, activation='relu', input_shape=(8, )))
model.add(keras.layers.Dense(1))

model.compile(
    optimizer='rmsprop',
    loss='mse',
    metrics=['mae']
)
```

```
In [21]: model.build()
```

```
In [22]: model.fit(features, labels, epochs=50, batch_size=32)
```

```
Epoch 1/50
639/639 [=====] - 2s 3ms/step - loss: 24784873472.0000 - mae: 123846.9688
Epoch 2/50
639/639 [=====] - 2s 3ms/step - loss: 24534173696.0000 - mae: 123053.7344
Epoch 3/50
639/639 [=====] - 2s 3ms/step - loss: 24286550016.0000 - mae: 122263.8125
Epoch 4/50
639/639 [=====] - 2s 3ms/step - loss: 24042905600.0000 - mae: 121491.8906
Epoch 5/50
639/639 [=====] - 2s 3ms/step - loss: 23798994944.0000 - mae: 120715.5703
Epoch 6/50
639/639 [=====] - 2s 3ms/step - loss: 23559567360.0000 - mae: 119947.5703
Epoch 7/50
639/639 [=====] - 2s 3ms/step - loss: 23321169920.0000 - mae: 119187.1641
Epoch 8/50
639/639 [=====] - 2s 3ms/step - loss: 23085172736.0000 - mae: 118433.7969
Epoch 9/50
639/639 [=====] - 2s 3ms/step - loss: 22850301952.0000 - mae: 117681.8359
Epoch 10/50
639/639 [=====] - 2s 3ms/step - loss: 22619633664.0000 - mae: 116939.9922
Epoch 11/50
639/639 [=====] - 2s 3ms/step - loss: 22389725184.0000 - mae: 116208.4141
Epoch 12/50
639/639 [=====] - 2s 3ms/step - loss: 22163886080.0000 - mae: 115477.0859
Epoch 13/50
639/639 [=====] - 2s 3ms/step - loss: 21939034112.0000 - mae: 114750.6562
Epoch 14/50
639/639 [=====] - 2s 3ms/step - loss: 21716807680.0000 - mae: 114036.6406
Epoch 15/50
639/639 [=====] - 2s 3ms/step - loss: 21498261504.0000 - mae: 113327.4922
Epoch 16/50
639/639 [=====] - 2s 3ms/step - loss: 21283094528.0000 - mae: 112626.4297
Epoch 17/50
639/639 [=====] - 2s 3ms/step - loss: 21069518848.0000 - mae: 111933.3750
Epoch 18/50
639/639 [=====] - 2s 3ms/step - loss: 20857874432.0000 - mae: 111247.3516
Epoch 19/50
639/639 [=====] - 2s 3ms/step - loss: 20650741760.0000 - mae: 110577.6016
Epoch 20/50
639/639 [=====] - 2s 3ms/step - loss: 20444252160.0000 - mae: 109907.5312
Epoch 21/50
639/639 [=====] - 2s 3ms/step - loss: 20243519488.0000 - mae: 109252.6875
```

```
Out[22]: <keras.callbacks.History at 0x7f5f7c1fd7d0>
```

```
In [35]: test_df = pd.DataFrame({
    'longitude': [-119.85],
    'latitude': [37.48],
    'housing_median_age': [22],
    'total_rooms': [2850],
    'total_bedrooms': [500],
    'population': [1150],
    'households': [460],
    'median_income': [3.12]
})
```

```
In [36]: test_df
```

```
Out[36]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-119.85	37.48	22	2850	500	1150	460	3.12

```
In [37]: test_df = test_df - mean
test_df = test_df / std
```

```
In [38]: test_df
```

```
Out[38]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-0.139406	0.864456	-0.526779	0.097698	-0.089872	-0.242627	-0.103148	-0.395496


```
In [39]: model.predict(test_df)
```

```
Out[39]: array([[38523.855]], dtype=float32)
```

The predicted median house value is 38523.855

```
In [40]:  
Out[40]: 0      452600.0  
          1      358500.0  
          2      352100.0  
          3      341300.0  
          4      342200.0  
          ...  
          20635    78100.0  
          20636    77100.0  
          20637    92300.0  
          20638    84700.0  
          20639    89400.0  
          Name: median_house_value, Length: 20433, dtype: float64
```

Learning Curve



Weights

```
In [154]: # values of the trained weights  
print(lin_reg.coef_)  
  
[ -418268.88736635  -387070.03453979   60557.81471414   77206.23571651  
   91860.60238544  -1469411.46051484   663328.64835277   618217.58800007  
   495057.21138978    71127.52066782   292806.6403382 ]
```

Training Loss Value:

```
In [161]: #final training loss value  
print('Your final training loss value rmse = {} R2 = {}'.format(final_rmse, r2))  
  
Your final training loss value rmse = {} R2 = {} 74292.19350137576 0.5788085928215299
```

Predicted Value for Test Case:

The predicted median house value is 38523.855

THE END